

CHRISTIAN BOURLIER

AI SYSTEMS ARCHITECT

christian@bourlier.ai 
310.963.5282 
christianbourlier 
christianbourlier.com 
rezzed.ai 

“

I Build Bounded AI Systems That Ship.

Observable · Vendor-neutral · Governed before execution · Designed for deployment

HOW I THINK

Systems First

Every problem is a system. I map inputs, outputs, feedback loops, constraints, and failure modes before code. Then I build the smallest architecture that satisfies them.

Governance Before Autonomy

Agents do not execute freely. They operate within identity boundaries, context budgets, rate limits, and policy gates — enforced before tool use.

Typed Work > Vanity Metrics

Not all tasks are equal. My orchestration differentiates WORK (sprint stories, builds, directives) from CONTROL (ACKs, handoffs, signals). Measurement clarity prevents metric inflation.

Compounding Infrastructure

I don't build isolated tools. I build composable layers — orchestration, safety, evaluation, dispatch — where each piece reinforces the system.

THE APPROACH

- Diagnose first.
- Constrain risk.
- Ship a bounded POC.
- Instrument reality.
- Iterate toward production.

THE CONVERGENCE

8+ years AI & data engineering

3,500+ typed autonomous tasks orchestrated

60M+ daily records processed (99.9% uptime)

\$10M+ ad spend optimized via privacy-preserving ML

“ Christian brings both technical excellence and great team spirit... a natural leader who motivates his team to accomplish their goals.

— Brianna Mersey, VP Data

CURRENT SYSTEMS

Rezzed.ai

AI-native orchestration studio

Multi-agent Claude Code architecture (Planner → Worker). CacheBash MCP relay. Guardian pre-tool-use safety gate. Vendor-neutral model routing. 25+ MCP tool definitions with schema validation. Open-source: dispatch, dreamwatch, dual-auth, guardian, poll-hooks, rate-gate, reach, specfirst, voicekeeper, xfetch.

[Multi-Agent](#) [MCP](#) [Safety](#) [Orchestration](#)

CacheBash

Async coordination for AI agents

Transactional task claiming. TTL-based relay messaging. Identity-scoped execution. Audit logging at tool boundary. Retry classification, provenance hashing, dynamic token budgeting.

[MCP Server](#) [Cloud Run](#) [Firestore](#)

OptiMeasure

Cookieless attribution — MMM-lite

Python + Vertex AI modeling. Privacy-first measurement. Reduced modeling latency from 6 months to 6 weeks.

[Vertex AI](#) [Privacy](#) [Python](#)

christianbourlier.com

Portfolio built as proof of execution

Next.js 16, Three.js GLSL, 8-palette engine, Zustand, PWA. Production Gemini integration with rate limiting, validation, audit logging.

[Next.js](#) [Three.js](#) [Gemini](#)

OPEN TO

- Applied AI / Forward Deployed roles
- Multi-agent architecture design
- AI governance & runtime safety systems
- Production LLM deployment
- Executive-technical bridge roles

Not just what AI can do — what it can safely and measurably deliver.

CERTS

- GCP Professional Data Engineer
- GCP Professional Cloud Architect

DIAGNOSE → CONSTRAIN → ORCHESTRATE → SHIP

IMPLEMENTATION

[Python](#) [SQL](#) [GCP](#) [BigQuery](#) [AWS](#) [Snowflake](#) [React](#) [Next.js](#) [Vertex AI](#)
[LLMs](#) [MCP](#) [Agentic Systems](#) [ETL/ELT](#) [Safety Gating](#) [Prompting](#) [Clean Rooms](#)
[Dashboarding](#) [Attribution](#) [Privacy](#) [GLSL](#) [Firestore](#) [Audit Logging](#)

GCP Professional Data Engineer · GCP Professional Cloud Architect