

# CHRISTIAN BOURLIER

AI SYSTEMS ARCHITECT

[github.com/rezzedai](https://github.com/rezzedai)

[christian@bourlier.ai](mailto:christian@bourlier.ai)

[bourlier.ai](https://bourlier.ai)

“ I build bounded AI systems that ship –

powerful enough to automate, constrained enough to trust.

## THE CONVERGENCE

**8+ years** AI & data engineering

**650+** typed tasks within bounded multi-agent orchestration

**3,500+** total tasks across orchestration and production systems

**60M+** daily records processed (99.9% uptime)

**\$10M+** ad spend optimized via privacy-preserving ML

“ Christian brings both technical excellence and great team spirit... a natural leader who motivates his team to accomplish their goals.

— Brianna Mersey, VP Data

## HOW I THINK

### Systems First

Every problem is a system. I map inputs, outputs, feedback loops, constraints, and failure modes before code. Then I build the smallest architecture that satisfies them.

### Typed Work > Vanity Metrics

My orchestration differentiates WORK (sprint stories, builds, directives) from CONTROL (ACKs, handoffs, signals). Measurement clarity prevents metric inflation.

### Constrained Autonomy

Agents operate within identity boundaries, context budgets, rate limits, and pre-execution policy gates. Orchestration, safety, evaluation, and dispatch layers reinforce bounded execution.

## THE APPROACH

- Diagnose.
- Bound.
- Orchestrate.
- Instrument.
- Deploy.
- Iterate.

### Instrumentation formalization underway:

- Explicit completion semantics (SUCCESS / FAILED / SKIPPED)
- Retry classification and error taxonomy
- Token, cost, and latency rollups per task class
- Prompt/output provenance hashing

## CURRENT SYSTEMS

### Rezzed.ai

*Bounded multi-agent orchestration*

Built on Claude Code. Coordinates role-specialized agents (Planner → Worker topology) within explicit execution envelopes using typed task lifecycles (WORK vs CONTROL), identity-scoped boundaries, context threshold monitoring, static token ceilings, and pre-tool-use policy gating. Model backends are adapter-based and vendor-neutral. Designed to enforce bounded autonomy rather than maximize agent freedom.

### CacheBash

*MCP relay with persistent task state*

Deployed on Cloud Run with transactional claiming. Implements TTL-based relay messaging with dead-letter containment, identity isolation across agents, and audit logging at the tool boundary. Differentiates productive work from control traffic to prevent metric inflation and clarify execution semantics. Orchestration logic remains decoupled from provider SDK assumptions.

### OptiMeasure

*Privacy-first AI measurement*

Python and Vertex AI architecture to accelerate cookieless attribution modeling. Reduced modeling latency from six months to six weeks while preserving clean-room compatibility and compliance constraints. Designed for transparent, deployable attribution under real budget and governance requirements.

### bourlier.ai

*Constraint-aware deployment system*

Production portfolio implemented as a live testbed rather than a static showcase. Integrates Gemini with rate limiting, validation, and audit logging; enforces controlled prompt design and structured output handling. Model integration patterns under bounded execution principles.

## INTERESTED IN

Applied AI systems where autonomy is governed by explicit boundaries and measurable outcomes — particularly orchestration, model deployment under real-world constraints, and evaluation-driven iteration.

## CERTS

- GCP Professional Data Engineer
- GCP Professional Cloud Architect