CHRISTIAN BOURLIER

AI SYSTEMS ARCHITECT

christian@bourlier.ai
310.963.5282
christianbourlier
christianbourlier.com
rezzed.ai

> *I Build Bounded AI Systems That Ship.*
>
> Observable · Vendor–neutral · Governed before execution · Designed for deployment

HOW I THINK

Systems First
Every problem is a system. I map inputs, outputs, feedback loops, constraints, and failure modes before code. Then I build the smallest architecture that satisfies them.

Governance Before Autonomy
Agents do not execute freely. They operate within identity boundaries, context budgets, rate limits, and policy gates — enforced before tool use.

Typed Work > Vanity Metrics
Not all tasks are equal. My orchestration differentiates WORK (sprint stories, builds, directives) from CONTROL (ACKs, handoffs, signals). Measurement clarity prevents metric inflation.

Compounding Infrastructure
I don't build isolated tools. I build composable layers — orchestration, safety, evaluation, dispatch — where each piece reinforces the system.

THE APPROACH

→ Diagnose first.
→ Constrain risk.
→ Ship a bounded POC.
→ Instrument reality.
→ Iterate toward production.

THE CONVERGENCE

8+ years   AI & data engineering

3,500+   typed autonomous tasks orchestrated

60M+   daily records processed (99.9% uptime)

$10M+   ad spend optimized via privacy-preserving ML

> *Christian brings both technical excellence and great team spirit... a natural leader who motivates his team to accomplish their goals.*
>
> — Brianna Mersey, VP Data

CURRENT SYSTEMS

Rezzed.ai
AI-native orchestration studio

Multi-agent Claude Code architecture (Planner → Worker). CacheBash MCP relay. Guardian pre-tool-use safety gate. Vendor-neutral model routing. 25+ MCP tool definitions with schema validation. Open-source: dispatch, dreamwatch, dual-auth, guardian, poll-hooks, rate-gate, reach, specfirst, voicekeeper, xfetch.

`Multi-Agent` `MCP` `Safety` `Orchestration`

CacheBash
Async coordination for AI agents

Transactional task claiming. TTL-based relay messaging. Identity-scoped execution. Audit logging at tool boundary. Retry classification, provenance hashing, dynamic token budgeting.

`MCP Server` `Cloud Run` `Firestore`

OptiMeasure
Cookieless attribution — MMM-lite

Python + Vertex AI modeling. Privacy-first measurement. Reduced modeling latency from 6 months to 6 weeks.

`Vertex AI` `Privacy` `Python`

christianbourlier.com
Portfolio built as proof of execution

Next.js 16, Three.js GLSL, 8-palette engine, Zustand, PWA. Production Gemini integration with rate limiting, validation, audit logging.

`Next.js` `Three.js` `Gemini`

OPEN TO

- Applied AI / Forward Deployed roles
- Multi-agent architecture design
- AI governance & runtime safety systems
- Production LLM deployment
- Executive-technical bridge roles

*Not just what AI can do —
what it can safely and measurably deliver.*

CERTS

- GCP Professional Data Engineer
- GCP Professional Cloud Architect

DIAGNOSE → CONSTRAIN → ORCHESTRATE → SHIP

IMPLEMENTATION

`Python` `SQL` `GCP` `BigQuery` `AWS` `Snowflake` `React` `Next.js` `Vertex AI`
`LLMs` `MCP` `Agentic Systems` `ETL/ELT` `Safety Gating` `Prompting` `Clean Rooms`
`Dashboarding` `Attribution` `Privacy` `GLSL` `Firestore` `Audit Logging`

GCP Professional Data Engineer · GCP Professional Cloud Architect

christianbourlier.com · rezzed.ai · github.com/rezzedai

Strength and Honour.